



r4 Inputs — DS Backlog Feeds into `r4 = _compute_backlog_balance(ds_backlog)`
(range $[-3, +1]$, 10 eval seeds averaged, 500-episode trained agent)

(1) DS 1 backlog over time



(2) Period-over-period change — `r4` penalizes $\Sigma(\Delta b) > 0$



(3) Realized `r4` over time — range $[-3, +1]$

