

Reward Components Over Time — DS 1 (DRL) Test Run
(10 eval seeds averaged, 500-episode trained agent, MAB off, equal weights, 95% severe disruption (thesis-faithful))

