



(1) Cumulative reward per episode (raw + smoothed)



(2) Mean DS backlog per episode — should decrease as agent learns



(3) Exploration rate decay

